


ramya ai

The Universal Hybrid Agent Platform That Slashes AI Costs by up to 80%

"We help computer users and professionals across all industries run powerful autonomous AI agents at a fraction of the cost by executing routine tasks locally on their hardware and offloading complex reasoning to optimized cloud models."



The Universal Token Burn

80%

of agentic AI queries are routine sub-tasks
like file search, text formatting and basic web search.

Who Faces This

Anyone using AI daily—developers, researchers, lawyers, accountants, students, and power users.

What Happens


Users are forced to pay frontier cloud rates for trivial everyday actions, rapidly exhausting monthly credits.

The Pain

Monthly software bills of \$20–\$50+/user, disruptive rate limits mid-task, and privacy concerns when uploading sensitive files.

Frequency

Experienced continuously, multiple times every single day.



Why Hasn't This Been Solved? The Missing Bridge

Users worldwide are forced to choose between two flawed extremes

Pure Cloud AI – ChatGPT / Claude

Powerful and easy to use, but expensive, has strict usage limits, and uploads all data to remote servers.

Pure Local Open-Source AI

Free and private, but requires complex technical setup and struggles on standard consumer hardware.

❏ **The Gap:** No simple desktop app bridges optimized local on-device execution with cost-efficient cloud escalation.

Today's Workaround: Juggling multiple paid subscriptions and copying data back and forth manually.

Solution: The ramya Hybrid Architecture

Universal, zero-friction desktop AI engine

1

1-Click Desktop App

Clean, zero-code interface designed for both technical and non-technical users.

2

Free & Private Local Engine

Executes file searches, text summaries, data formatting, and routine sub-tasks directly on your device.

3

Cost-Optimized Cloud Mesh

Seamlessly offloads complex multi-step reasoning to fine-tuned high-intelligence cloud models.

4

User Choice & Flexibility

Local inference is 100% free; users pay only for cloud power at up to 80% lower cost, with Bring Your Own Key (BYOK) support.



Our First Beachhead Customer

PERSONA

High-Frequency AI Power Users & Knowledge Professionals

Who They Are

Active freelancers, developers, consultants, and researchers using AI 3+ hours daily.

Their Situation

Deeply reliant on AI for writing, coding, data analysis, and document management.

Their Pain

Spending \$50–\$150/month across fragmented tools and constantly hitting message caps.

Current Alternative

Juggling simultaneous subscriptions (ChatGPT Plus, Claude Pro, Cursor).

Why They Pay

One unified desktop app that saves them 70%+ on monthly AI expenses and eliminates rate limits.

Revenue Model

\$20/month market standard with 3x the capacity.

Free (\$0)

- Unlimited local on-device AI
- Free 500k starter cloud tokens
- \$0 server cost per user

Ramya Pro (\$20/mo)

- **3x token capacity vs. ChatGPT**
- Zero rate limits on heavy tasks
- Cross-device workflow sync

Pay-As-You-Go & API

- Instant \$5 / \$10 token reload packs
- Developer API access (per 1M tokens)
- Team & enterprise seat scaling

Traction, Team & Next Steps

Validation & Progress

- Built functional open-source AI project 'Jarvis' on GitHub.
- Fine-tuned lightweight models on Hugging Face for token-efficient tool-calling.

Team & Capability

- Hands-on technical background in Transformer architectures, Reinforcement Learning, and local quantization.
- Deep dedication to eliminating the hardware and financial barriers of AI agents.

The Ask

- Seeking incubation support, cloud compute credits, and early beta testers for our global desktop release.